Guidelines

Jiafeng Chen   Yufeng Ling   Francisco Rivera

Abstract

- This document describes the expected style, structure, and rough proportions for your final project write-up.
- While you are free to break from this structure, consider it a strong prior for our expectations of the final report.
- Length is a hard constraint. You are only allowed max **8 pages** in this format. While you can include supplementary material, it will not be factored into the grading process. It is your responsibility to convey the main contributions of the work in the length given.

1. Introduction

Example Structure:

- What is the problem of interest and what (high-level) are the current best methods for solving it?

- How do you plan to improve/understand/modify this or related methods?

- Preview your research process, list the contributions you made, and summarize your experimental findings.

2. Background

Example Structure:

- What information does a non-expert need to know about the problem domain?

- What data exists for this problem?

- What are the challenges/opportunities inherent to the data? (High dimensional, sparse, missing data, noise, structure, discrete/continuous, etc?)

3. Related Work

Example Structure:

- What 3-5 papers have been published in this space?

- How do these differ from your approach?

- What data or methodologies do each of these works use?

- How do you plan to compare to these methods?

4. Model

Example Structure:

- What is the formal definition of your problem?

- What is the precise mathematical model you are using to represent it? In almost all cases this will use the probabilistic language from class, e.g.

$$z \sim \mathcal{N}(0, \sigma^2) \tag{1}$$

But it may also be a neural network, or a non-probabilistic loss,

$$h_t \leftarrow \mathrm{RNN}(x_t, h_{t-1})$$

This is also a good place to reference a diagram such as Figure 1.

- What are the parameters or latent variables of this model that you plan on estimating or inferring? Be explicit. How many are there? Which are you assuming are given? How do these relate to the original problem description?

5. Inference (or Training)

- How do you plan on training your parameters / inferring the states of your latent variables (MLE / MAP / Backprop / VI / EM / BP / ...)

- What are the assumptions implicit in this technique? Is it an approximation or exact? If it is an approximation what bound does it optimize?

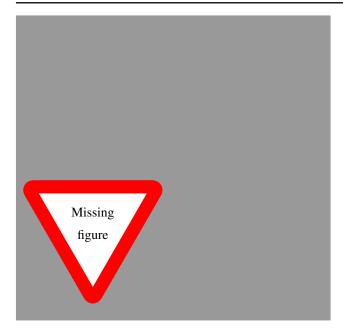- What is the explicit method / algorithm that you derive for learning these parameters?

Figure 1. This is a good place to include a diagram showing how your model works. Examples include a graphical model or a neural network block diagram.
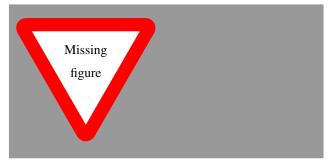

Algorithm 1 Your Pseudocode

---

6. Methods

- What are the exact details of the dataset that you used? (Number of data points / standard or non-standard / synthetic or real / exact form of the data)

- What are the exact details of the features you computed?

- How did you train or run inference? (Optimization method / hyperparameter settings / amount of time ran / what did you implement versus borrow / how were baselines computed).

- What are the exact details of the metric used?

7. Results

- What were the results comparing previous work / baseline systems / your systems on the main task?

- What were the secondary results comparing the variants of your system?

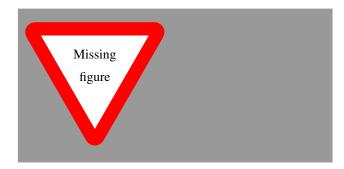- This section should be fact based and relatively dry. What happened, what was significant?

Figure 2. Visualizations of the internals of the system.

8. Discussion

- What conclusions can you draw from the results section?

- Is there further analysis you can do into the results of the system? Here is a good place to include visualizations, graphs, qualitative analysis of your results.

- What questions remain open? What did you think might work, but did not?

9. Conclusion

- What happened?

- What next?

Table 1. This is usually a table. Tables with numbers are generally easier to read than graphs, so prefer when possible.
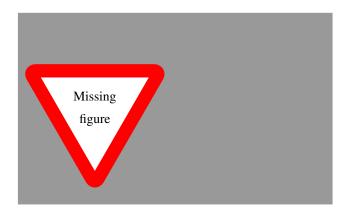

Table 2. Secondary table or figure in results section.